

# Statistical Methods in Natural Language Processing (NLP)



*Class 5: Introduction to Python: Pandas and Matplotlib*

Charalambos (Haris) Themistocleous

*Department of Philosophy, Linguistics and Theory  
of Science, Centre for Linguistic Theory and Studies  
in Probability*

# Introduction to Python

1. Series and DataFrames in Python
2. Basic Plots
3. Descriptive Statistics in Python
4. Distributions in SciPy

# Introduction to Python

1. Python 2.7 or higher (including Python 3)
2. pandas
3. NumPy
4. matplotlib
5. IPython
6. NLTK

# Pandas, Matplotlib, NLTK

- ▶ **Pandas**, is a package that provides functionality for analyzing data in the form of tables, such as those we have in Excel, Libreoffice Calc. The most important data structure is the DataFrame which is very similar to R dataframes. Pandas also provide functionality for reshaping, sorting, manipulating, etc., data.
- ▶ The second library we will be using is **NumPy**, which offers the basic functionality for conducting mathematics, including statistics, linear algebra, and Fourier transformations.
- ▶ **Matplotlib** provides functionality for creating plots and graphs.
- ▶ **NLTK** is a Natural Language Toolkit implemented in Python.
- ▶ So, to start an analysis add the following code on your code file. The code imports the libraries and provide a designated name for each library. So, we will be calling pandas for instance we will use the name `pd` followed by a period and the name of a function. This will become more clear soon.

# Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import nltk
```

# Data Manipulation with Pandas

1. Series
2. DataFrames
3. Date-Time objects

# Series

A Series is a single vector of data with an index for each element. A similar structure in numpy is the array.

```
measurements = pd.Series([328259, 22781, 30857, 4164, 328387])  
measurements
```

The printed output is the following:

```
Out[1]:  
0      328259  
1       22781  
2       30857  
3         4164  
4      328387  
dtype: int64
```

# Values and Indexes

Series consist of **values** and **indexes**, we can call them separately in the following manner:

```
measurements.values
```

```
Out[2]:
```

```
array([328259,  22781,  30857,   4164, 328387])
```

```
measurements.index
```



# Selecting Values in Series

```
measurements[3]  
Out[4]:  
4164
```

## Selecting Values in Series

We can select values based on logical operations as well

```
measurements[measurements < 20000]
```

```
Out[5]:
```

```
Ecuador      4164
```

```
dtype: int64
```

or

```
measurements[measurements == 22781]
```

```
Out[6]:
```

```
Argentina      22781
```

```
dtype: int64
```

## Providing optional labels to Series

These numbers are not very informative so we want to provide labels. So, if we know that these numbers represent the number of books published in 2010 we might want to provide the name of the country as an index.

```
measurements = pd.Series([328259, 22781, 30857, 4164, 328387],  
                          index=['USA', 'Argentina', 'Sweden', 'Ecuador', 'China'])
```

```
measurements
```

```
Out[7]:
```

USA	328259
Argentina	22781
Sweden	30857
Ecuador	4164
China	328387

```
dtype: int64
```

## Selecting Values in Series using labels

We can use these labels to select the value.

```
measurements [ 'USA ' ]  
328259
```

## Series and labels

Also, we can provide labels both to the array of values and to the index:

```
measurements.name = 'Book Counts'
measurements.index.name = 'Countries'
measurements
Out[71]:
Countries
USA          328259
Argentina    22781
Sweden       30857
Ecuador      4164
China        328387
Name: Book Counts, dtype: int64
```

## Selecting Values in Series using specific criteria

We might be interested to select only the countries whose name ends in letter 'a':

```
measurements[[name.endswith('a') or name.endswith('A') for name in measurements.index]]  
USA          328259  
Argentina    22781  
China        328387  
dtype: int64
```

The following provides information about the position of these numbers:

```
[name.endswith('a') or name.endswith('A') for name in measurements.index]  
[True, True, False, False, True]
```

# Maths and Series

NumPy's math functions and statistics can be applied to Series, e.g.,

```
np.mean(measurements)  
142889.6
```

## Series and Dictionaries

Series are very similar objects to standard dictionaries (dict) in Python:

```
Bookpublications = {'Italy':59743, 'Argentina':22781,  
                    'Poland': 31500, 'Vietnam': 24589, 'Indonesia': 24000}
```

```
pd.Series(Bookpublications)
```

```
Argentina    22781  
Indonesia    24000  
Italy        59743  
Poland       31500  
Vietnam      24589  
dtype: int64
```





# DataFrame

```
data = pd.DataFrame(  
{ 'counts': [ 328259, 22781, 30857, 4164, 328387, 59743, 31500, 24589 ],  
  'year': [2010, 2010, 2010, 2010, 2010, 2005, 2010, 2009],  
  'country': [ 'USA', 'Argentina', 'Sweden', 'Ecuador',  
               'China', 'Italy', 'Poland', 'Vietnam' ] } )  
data
```

# DataFrame

The output now is a table as we expect it to be:

	country	counts	year
0	USA	328259	2010
1	Argentina	22781	2010
2	Sweden	30857	2010
3	Ecuador	4164	2010
4	China	328387	2010
5	Italy	59743	2005
6	Poland	31500	2010
7	Vietnam	24589	2009

## Selecting Values from DataFrames

To select the values of the column, we can use its name:

```
data[ ' counts ' ]
```

```
0      328259
1       22781
2       30857
3        4164
4      328387
5       59743
6       31500
7       24589
```

```
Name: counts, dtype: int64
```

## Selecting Values from DataFrames

```
data.counts
```

```
Out[95]:
```

```
0      328259
```

```
1       22781
```

```
2       30857
```

```
3        4164
```

```
4      328387
```

```
5       59743
```

```
6       31500
```

```
7       24589
```

```
Name: counts, dtype: int64
```



## Changing the order of columns

```
data[['country', 'year', 'counts']]
```

# Indexes

The index of columns is provided by the following:

```
data.columns
```

```
Out[91]:
```

```
Index(['country ', 'counts ', 'year '], dtype='object ')
```

# Types and selections

```
type(data.counts)  
pandas.core.series.Series
```

```
type(data[['counts']])  
pandas.core.frame.DataFrame
```

## Selecting Rows

To select a row in a DataFrame, we index its *ix* attribute in the following way:

```
data.ix[3]  
Out[98]:  
country      Ecuador  
counts        4164  
year          2010  
Name: 3, dtype: object
```



## Dictionaries and DataFrames

We might create DataFrames using dictionaries

Alternatively , we can create a DataFrame with a dict of dicts:  
In [111]:

```
data = pd.DataFrame(  
    {0:{ 'AA': 1, 'gender': 'Male', 'height': 168},  
    1: { 'AA': 2, 'gender': 'Male', 'height': 180},  
    2: { 'AA': 3, 'gender': 'Female', 'height': 170},  
    3: { 'AA': 4, 'gender': 'Female', 'height': 169},  
    4: { 'AA': 5, 'gender': 'Female', 'height': 170},  
    5: { 'AA': 6, 'gender': 'Male', 'height': 165}})
```

In [112]:

data

Out[112]:

	0	1	2	3	4	5
AA	1	2	3	4	5	6
gender	Male	Male	Female	Female	Female	Male
height	168	180	170	169	170	165

# Transpose Function

To get the 'standard' DataFrame output we need to transpose the code:

```
data = data.T
```

```
data
```

```
Out[113]:
```

	AA	gender	height
0	1	Male	168
1	2	Male	180
2	3	Female	170
3	4	Female	169
4	5	Female	170
5	6	Male	165

## Indexes and values

DataFrames have indexes and values which are called in the following way:

```
data.values
```

The output is following

```
array([[1, 2, 3, 4, 5, 6],  
       ['Male', 'Male', 'Female', 'Female', 'Female', 'Male'],  
       [168, 180, 170, 169, 170, 165]], dtype=object)
```

# Indexes and values

and the index is called by `data.index` and the result is:

```
Index(['AA', 'gender', 'height'], dtype='object')
```

We cannot change the index, if we try, e.g., `data.index[1] = 5` Python will provide the following message: “Index does not support mutable operations”.

## Selecting columns and changing values

To select a column:

```
heights = data.height  
heights  
Out[116]:  
0      168  
1      180  
2      170  
3      169  
4      170  
5      165  
Name: height, dtype: object
```

## Selecting columns and changing values

To change a value

```
heights[5] = 191
```

```
heights
```

```
Out[117]:
```

```
0    168
```

```
1    180
```

```
2    170
```

```
3    169
```

```
4    170
```

```
5    191
```

```
Name: height, dtype: object
```

```
data
```

```
Out[118]:
```

```
AA  gender  height
```

```
0  1    Male    168
```

```
1  2    Male    180
```

```
2  3  Female    170
```

```
3  4  Female    169
```

```
4  5  Female    170
```

```
5  6    Male    191
```

# The copy function

```
ht = data.height.copy()  
ht[5] = 180
```

```
data
```

```
Out[141]:
```

	AA	gender	height
0	1	Male	168
1	2	Male	180
2	3	Female	177
3	4	Female	169
4	5	Female	170
5	6	Male	191

## Create/ modify columns by assignment

```
data.height[2] = 177
```

```
data
```

```
Out[122]:
```

	AA	gender	height
0	1	Male	168
1	2	Male	180
2	3	Female	177
3	4	Female	169
4	5	Female	170
5	6	Male	180



## Create/ modify columns by assignment

```
data['Status'] = 'Printed'
```

```
data
```

```
Out[143]:
```

	AA	gender	height	Status
0	1	Male	168	Printed
1	2	Male	180	Printed
2	3	Female	177	Printed
3	4	Female	169	Printed
4	5	Female	170	Printed
5	6	Male	191	Printed

## Create/ modify columns by assignment

The following method does not create a column:

```
data.libraryNo = 999
```

```
data
```

```
Out[146]:
```

	AA	gender	height	Status
0	1	Male	168	Printed
1	2	Male	180	Printed
2	3	Female	177	Printed
3	4	Female	169	Printed
4	5	Female	170	Printed
5	6	Male	191	Printed

## Create/ modify columns by assignment

```
data.libraryNo  
999
```

# Create DataFrame Columns using Series

We can define a Series object as column in a DataFrame

```
test = pd.Series([0]*2 + [3]*2)  
test
```

```
data['test'] = test  
data
```

# Strings in DataFrames

We created a Series of 4 numbers. Note however that the DataFrame contains six rows. This is not a problem when we use numbers because Python automatically add NaN to fill the empty rows. Nevertheless, when we employ other data structures such as strings Python will show an error message: `ValueError: Length of values does not match length of index.`

```
# Popular Authors
authors = ['Stephen King', 'J.K. Rowling', 'Mark Twain', 'George R. R. Martin']
data['authors'] = authors
```

## Strings in DataFrames

To correct the error, we simply add a string Series that has the same length as the DataFrame

```
authors = ['Stephen King', 'J.K. Rowling', 'Mark Twain',
           'George R. R. Martin', 'Charles Dickens', 'Arthur Conan Doyle']
data['favorite_authors'] = authors
```

This time the output is correct:

	AA	gender	height	Status	test	favorite_authors
0	1	Male	168	Printed	0.0	Stephen King
1	2	Male	180	Printed	0.0	J.K. Rowling
2	3	Female	177	Printed	3.0	Mark Twain
3	4	Female	169	Printed	3.0	George R. R. Martin
4	5	Female	170	Printed	NaN	Charles Dickens
5	6	Male	165	Printed	NaN	Arthur Conan Doyle

## Deleting Columns in DataFrames

To delete the column test from the DataFrame data

```
del data['test ']
```

```
data
```

	AA	gender	height	Status	authors
0	1	Male	168	Printed	Stephen King
1	2	Male	180	Printed	J.K. Rowling
2	3	Female	177	Printed	Mark Twain
3	4	Female	169	Printed	George R. R. Martin
4	5	Female	170	Printed	Charles Dickens
5	6	Male	165	Printed	Arthur Conan Doyle

## DataFrame as a simple ndarray

To get the data as a simple ndarray we need to employ the attribute `values`.

```
array([[1, 'Male', 168, 'Printed', 'Stephen King'],  
       [2, 'Male', 180, 'Printed', 'J.K. Rowling'],  
       [3, 'Female', 177, 'Printed', 'Mark Twain'],  
       [4, 'Female', 169, 'Printed', 'George R. R. Martin'],  
       [5, 'Female', 170, 'Printed', 'Charles Dickens'],  
       [6, 'Male', 165, 'Printed', 'Arthur Conan Doyle']],  
      dtype=object)
```

The `dtype` here is “object” because we have numeric and string data and differs when we have numeric or other type of data.



# Merging DataFrames

```
df1 = pd.DataFrame('A': ['A0',  
    'A1', 'A2', 'A3'], 'B': ['B0', 'B1',  
    'B2', 'B3'], 'C': ['C0', 'C1', 'C2',  
    'C3'], 'D': ['D0', 'D1', 'D2', 'D3'],  
    index=[0, 1, 2, 3])
```

Example from

<http://pandas.pydata.org/pandas-docs/stable/merging.html>

df1				
	A	B	C	D
0	A0	B0	C0	D0
1	A1	B1	C1	D1
2	A2	B2	C2	D2
3	A3	B3	C3	D3

df2				
	A	B	C	D
4	A4	B4	C4	D4
5	A5	B5	C5	D5
6	A6	B6	C6	D6
7	A7	B7	C7	D7

df3				
	A	B	C	D
8	A8	B8	C8	D8
9	A9	B9	C9	D9
10	A10	B10	C10	D10
11	A11	B11	C11	D11

Result				
	A	B	C	D
0	A0	B0	C0	D0
1	A1	B1	C1	D1
2	A2	B2	C2	D2
3	A3	B3	C3	D3
4	A4	B4	C4	D4
5	A5	B5	C5	D5
6	A6	B6	C6	D6
7	A7	B7	C7	D7
8	A8	B8	C8	D8
9	A9	B9	C9	D9
10	A10	B10	C10	D10
11	A11	B11	C11	D11

# Date and Time

Python can manipulate date and time objects using the `datetime` module. It allows the production of calculations using time and date objects and also provides classes for controlling the output (see also, <https://docs.python.org/2/library/datetime.html>)

```
from datetime import datetime
#00%
now = datetime.now()
now
```

and the result is

```
datetime.datetime(2017, 1, 6, 14, 41, 4, 481168)
```

# Date and Time

To get the date only

```
#%%  
now.date()
```

and the output in this case is `datetime.date(2017, 1, 6)`. To find the day

```
#%%  
now.day
```

and the output is 6.

# Date and Time

Also, for the time

```
#0%  
now.time()
```

and the output is `datetime.time(14, 41, 4, 481168)`. We can also ask which is the week day:

```
#0%  
now.weekday()
```

that will generate the output 4

```
#0%  
from datetime import date, time
```

# Date and Time

```
#%%  
time(3, 24)
```

```
#%%  
age = now - datetime(1980, 8, 16)  
age/365
```

```
#%%  
days=(datetime(2017, 3, 10) - datetime(2017, 8, 16))  
days.days
```

# Importing data

We suggest that you use comma-separated value or CSV files when interacting with Python and other statistical software. In computing, CSV files store tabular data (numbers and text) in plain text. Columns are separated by commas; rows are terminated by newlines. This file format is not proprietary, the files can be edited in text editors and spreadsheet software, such as Excel and Calc.

# Importing data

```
dur = pd.read_csv("data/duration.csv", sep=";")
```

```
dur
```

```
Out[153]:
```

	experiment	duration
0	A	199
1	A	184
2	A	242
3	A	236
4	A	216
5	A	176
6	A	223
7	A	186
8	A	210
9	A	220
..	...	...
95	C	221
96	C	239
97	C	235
98	C	248
99	C	204
100	C	226
101	C	206
102	C	194
103	C	205
104	C	182

```
[105 rows x 2 columns]
```

# Importing data

We can also import another dataframe and add a column titled AA.

```
fricative = pd.read_table("data/fricatives.csv", sep=',')  
fricative['AA'] = pd.Series(range(1,8827))
```



# Head

```
# %%
```

```
fricative.head()
```

```
# %%
```

	id	duration	intensity	cog	sdev	skew	kurt	\
0	1	0.060398	32.671794	757.605236	1104.704765	13.835014	210.523631	
1	2	0.045656	38.906220	732.582945	1065.089424	12.654465	186.856393	
2	3	0.050907	47.209304	647.696728	1627.357767	7.647966	61.615315	
3	4	0.051049	41.703970	1017.179353	2318.797907	5.570367	33.783925	
4	5	0.028408	44.345609	1132.524942	848.894793	7.105495	108.453910	

	Segment	Vowel	Variety	Stress	Voice	Position	AA
0	d	a	CG	Unstressed	Voiced	Middle	1.0
1	d	a	CG	Unstressed	Voiced	Middle	2.0
2	d	a	CG	Unstressed	Voiced	Middle	3.0
3	d	a	CG	Unstressed	Voiced	Middle	4.0
4	d	a	CG	Unstressed	Voiced	Middle	5.0

## Skipping Rows

We can skip rows if we do not want them in the analysis:

```
# %%  
testfric=pd.read_csv("data/fricatives.csv", skiprows=[2,3,4,5,6])  
len(testfric.index)
```

To import a small number of rows from, we can use `nrows`:

```
# %%  
pd.read_csv("data/fricatives.csv", nrows=4)
```

## Skipping Rows

```
# %%  
pd.read_csv("data/fricatives.csv").head(20)  
  
pd.isnull(pd.read_csv("data/fricatives.csv")).head(20)
```

# Empty Cells

When we import data Python identifies empty cells, or NA values as NA data; to designate that specific values or symbols should be considered NA values, we can specify this as follows

```
pd.read_csv("data/fricatives.csv",  
na_values=['?', -9999999]).head(20)
```

# Saving Data

There are different methods to save data. To save data in CSV format

```
# ## Writing Data to Files  
fricative.to_csv("fricative-01.csv")
```

# Creating Plots using Pandas


Using pandas we can also make some basic plotting.

```
fricative['duration'].plot()
```

# Creating Plots using Pandas


```
#####
```

```
# %%  
fricative = pd.read_csv("data/fricatives.csv", sep=',')  
fricative['duration'].plot()  
# %%  
fricative['duration'].plot(kind='hist')  
  
# %%  
fricative['duration'].plot(kind='box', showfliers=False)
```






## Creating Plots using Pandas



## Creating Plots using Pandas



# Creating Plots using Matplotlib

```
import matplotlib.pyplot as plt  
plt.plot([1,2,3,4])  
plt.ylabel('some numbers')  
plt.show()
```

# Basic Descriptive Statistics using Pandas

In [208]:

```
fricative.sum()
```

Out[208]:

duration	827.811
intensity	346024
cog	5.05981e+07
sdev	2.40776e+07
skew	21699.9
kurt	328392
Segment	dddddddddddddddddddddddddddddddddddddd ...
Vowel	aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa ...
Variety	CGCGCGCGCGCGCGCGCGCGCGCGCGCGCGCGCG...
Stress	UnstressedUnstressedUnstressedUnstressedUnstre ...
Voice	VoicedVoicedVoicedVoicedVoicedVoicedVoicedVoic ...
Position	MiddleMiddleMiddleMiddleMiddleMiddleMiddleMiddleMid ...
AA	3.89536e+07

dtype: object

# Calculating the Mean

In [209]:

```
fricative.mean()
```

Out[209]:

```
duration      0.093782
intensity     39.254023
cog           5732.200660
sdev          2727.724598
skew          2.458354
kurt          37.203150
AA            4413.500000
dtype: float64
```

# Calculating the Standard Deviation

```
In [211]:
```

```
fricative.std()
```

```
Out[211]:
```

```
duration      0.031759
intensity     8.272744
cog           3425.508087
sdev          1339.636724
skew          4.785687
kurt          138.622132
AA            2547.991071
dtype: float64
```

# Descriptive Statistics and Counts

In [212]:

```
fricative.count()
```

Out[212]:

duration	8827
intensity	8815
cog	8827
sdev	8827
skew	8827
kurt	8827
Segment	8827
Vowel	8827
Variety	8827
Stress	8827
Voice	8827
Position	8827
AA	8826

dtype: int64

## Finding Missing Values and NaNs

```
fricative.intensity.hasnans
```

```
Out[215]:
```

```
True
```

```
In [221]:
```

```
fricative.intensity.isnull().sum()
```

```
Out[221]:
```

```
12
```



# Descriptive Statistics: Describe Function

Describe:

In [222]:

```
fricative.describe()
```

Out[222]:

	duration	intensity	cog	sdev	skew \
count	8827.000000	8815.000000	8827.000000	8827.000000	8827.000000
mean	0.093782	39.254023	5732.200660	2727.724598	2.458354
std	0.031759	8.272744	3425.508087	1339.636724	4.785687
min	0.020333	5.278827	419.757883	228.697624	-5.250996
25%	0.071596	NaN	2385.869561	1771.421219	-0.113557
50%	0.091452	NaN	6175.724355	2368.203536	0.925865
75%	0.112412	NaN	8344.008050	3595.757817	2.953676
max	0.316844	69.455969	18606.542539	9253.436646	59.853567

	kurt	AA
count	8827.000000	8826.000000
mean	37.203150	4413.500000
std	138.622132	2547.991071
min	-1.892874	1.000000
25%	0.512395	NaN
50%	3.432032	NaN
75%	12.453753	NaN
max	3999.613892	8826.000000

describe can detect non-numeric data and sometimes yield useful information about it.

# Descriptive Statistics: Describe Function

```
fricative.sdev.describe()  
Out[224]:  
count      8827.000000  
mean       2727.724598  
std        1339.636724  
min         228.697624  
25%        1771.421219  
50%        2368.203536  
75%        3595.757817  
max         9253.436646  
Name: sdev, dtype: float64
```

# Pantas and Scipy

- ▶ Scipy provides mathematical functions
- ▶ For more information see  
<http://docs.scipy.org/doc/scipy/reference/stats.html>

# Probability distributions in Python

- ▶ make a random variable representing a dice (unif distribution)

```
from scipy.stats import randint  
dice = randint(1, 7)
```

- ▶ roll the dice 1000 times


```
outcomes = dice.rvs(1000)
```

- ▶ What is the probability of getting 5?

```
dice.pmf(5)
```

- ▶ You can also count the mean(), variance, and standard deviation:  
dice.mean(), dice.var(), dice.std().

# Plotting the Distribution in Python



```
import scipy.stats
from matplotlib import pyplot as plt
dice = scipy.stats.randint(1, 7)
rolls = [1,2,3,4,5,6]
pmf_values = dice.pmf(rolls)
plt.bar(rolls, pmf_values, width=0.2)
# some cosmetics
plt.axis([0, 7, 0, 1])
plt.xlabel('Rolls')
plt.ylabel('Probability')
plt.grid()
plt.show()
# plt.savefig('dice_rolls.png')
```

Figure: Probabilities of dice rolls.

# Next Class

- ▶ Hypothesis Testing
- ▶ Statistical Models